



A GUIDE TO SQLITE_ORM FOR SQL AND C++ USERS

By Juan Dent © 2022 Version 2.3

Table of Contents

C++ for writing data intensive applications	6
A DSL has the following properties:	6
Working level	7
Mapping types to tables – making types persistent.....	8
Simple query	8
Simple calculation:	8
General SELECT Syntax:	9
One table select:	9
When dealing with large resultsets	10
Sorting rows.....	10
Order by	10
General syntax:	10
Simple order by:.....	10
Compound order by:.....	10
Dynamic Order by	11
Ordering by a function	11
Dealing with NULLs when sorting	11
Sorting on a data dependent key.....	12
Filtering Data	12
Select distinct	12
Where Clause	12
Limit.....	13
Between Operator.....	14
In	15
Like	16

Glob	17
IS NULL	17
Dealing with NULL values in columns.....	17
Joining tables	18
SQLite join	18
Cross join	18
Other joins.....	19
Joins having to do with NULLs (i.e. OUTER JOINS)	20
Example for Left Join:.....	22
Example for Inner Join.....	22
Example for Natural Join	23
Self join.....	24
Full outer join	25
Grouping data.....	27
Group by.....	27
Having.....	28
Set operators	28
Union	28
Stacking one resultset on top of another	31
Except.....	31
Intersect	32
Subquery.....	33
Subquery	33
Exists.....	35
More querying techniques	37
Case	37
Aliases for columns and tables.....	38

Applying aliases to CASE.....	40
Changing data	40
Inserting a single row into a table	40
Inserting an object	41
Inserting several rows	41
Inserting several objects via containers	42
Inserting several rows (becomes an update if primary key already exists)	43
Inserting only certain columns (provided the rest have either default_values, are nullable, are autoincrement or are generated):	43
Inserting from select – getting rowid (since primary key is integral)	44
Inserting default values:.....	44
Non-standard extension in SQLITE	44
Update.....	45
Update several rows.....	46
Update one row	46
Update an object.....	46
Delete Syntax	47
Delete rows that satisfy a condition.....	47
Delete all objects of a certain type.....	47
Delete a certain object by giving its primary key	47
Replace	47
Transactions.....	48
Transactions	48
Core functions	50
Data definition	54
Sqlite data types.....	54
Create table	55
CHECK constraint.....	56

Columns with specific collation and tables with Primary Key	57
FOREIGN KEY constraint	57
AUTOINCREMENT property	59
GENERATED COLUMNS	60
Databases may be created in memory if desired	61
INDEX and UNIQUE INDEX	61
DEFAULT VALUE for DATE columns	62
PERSISTENT collections	63
GETTERS and SETTERS	64
DEFINING THE SHEMA FOR SELF-JOINS	65
SUBENTITIES	66
UNIQUENESS AT THE COLUMN AND TABLE LEVEL	68
NOT NULL CONSTRAINT	68
VACUUM	68
Why do we need vacuum?	68
Triggers	69
What is a Trigger?	69
Why do we need them?	69
Syntax:	69
Accessing old and new column values according to action	69
Examples of Triggers	70
Data migration	73
About correct order of dropping/loading tables	75
Example	75
How to drop data without losing it	76
Interpret values in non-nullable column as nullable	77
Making a backup of the entire database	77

Ensuring that a column contains unique values before making the column unique	78
SQLite tools.....	78
SQLiteStudio and Sqlite3 command shell.....	78
Installing SQLite using vcpkg	79
SQLite import and export CSV.....	79
SQLite resources	79
Debugging tips	80
The Future of sqlite_orm	81
References	81

C++ for writing data intensive applications

C++ is a large language with a very expressive and rich syntax. Perhaps its most salient characteristic is its ability to control the level of abstraction enabling application programming to be done in terms of the problem domain's concepts. An extension to this capability is the ability to define Domain Specific Languages (DSL) like SQL in terms of "ordinary" C++ code. This makes C++ a very compelling language for writing data intensive applications.

A DSL has the following properties:

- It is a language, that defines: [CPPTMP,216]
 - An alphabet (set of symbols)
 - Well defined rules stating how the alphabet may be used to build well-formed compositions
 - A well-defined subset of all well-formed compositions that are assigned specific meanings
- It is domain specific not general-purpose
 - Examples include regular expressions, UML, Morse code
 - By this restriction, we gain significantly higher level of abstraction and expressiveness because
 - Specialized alphabet and notations allow pattern-matching that matches our mental model
 - Enabling writing code in terms close to the abstractions of the problem domain is the characteristic property and motivation behind all DSLs
 - We use the language's notation to write down a statement of the problem itself and the language's semantics take care of generating a solution
 - The most successful DSLs are often **declarative** languages providing us with notations to describe what rather than how
 - The how can be seen as a consequence of the what
- In a sense, DSL is an enhancement to object-oriented programming in which the development is done in terms of the problem domain conceptual model
 - We are just taking an extra step towards enriched notational support

This document refers to a SQL DSL called ***SQLITE ORM***, which provides direct support for writing SQL in C++. This is indeed a worthy capability and one that allows for the clear concise creation of data intensive applications. This library is not only a SQL DSL but a sort of object relational tool (ORM¹) in that it provides means to associate data structures in C++ with relational tables in sqlite3.

One can truly raise C++ abstraction level by thinking in a combination of imperative C++ enhanced with compile time metaprogramming and SQL. The synergy is indeed attractive and powerful.

¹ Object Relational Mapper: ORM

Working level

SQLITE_ORM allows us to interact with the persistent objects in two fundamentally different styles:

- By columns (pure SQL)
- By objects (SQL mapped to structures)

It is this support for dealing with persistency at the object level that explains the ORM suffix of the library name. Basically, each normalized table in the database may be represented in an application as a struct or class in a 1-to-1 relationship. This allows us to work at a high level of abstraction. We refer to these instances as “persistent atoms” to indicate their normalization and their unbreakable nature.

On the other hand, we have access to all the power and expression of SQL by allowing us to define queries for reading or writing in terms of columns just like you would if working in a relational client, but by virtue of this library being a Domain Specific Language (DSL)², it allows us to use C++ code to write SQL. This document is dedicated to every user of the library and can be thought of a dictionary of sorts between SQL and SQLITE_ORM.

The object queries deal with collections of instances of the persistent table associated with the object type. The column queries accessing persistent tables, deal at the column level of the corresponding object types.

² See Chapter 10 of C++ Template Metaprogramming, by David Abrahams and Aleksey Gurtovoy

Mapping types to tables – making types persistent

In order to work with persistent types we need to map them to tables and columns with a call to `make_storage(...)` like in this example for type `User`:

```
struct User {
    int id = 0;
    std::string name;
};
auto storage = make_storage(
    {dbName},
    make_table("users",
        make_column("id", &User::id, primary_key()),
        make_column("name", &User::name));
storage.sync_schema(); // synchronizes memory schema (called storage) with database schema
```

This creates non-nullable columns by default. To make one column nullable, say `name`, we must declare `name` to have one of these types:

- `std::optional<std::string>`
- `std::unique_ptr<std::string>`
- `std::shared_ptr<std::string>`

Another point to have in mind is that the fields may have any bindable type which includes all fundamental C++ data types, `std::string` and `std::vector`. Other types can be used but you must provide some code to make them bindable (an example is `std::chrono::sys_days`). In particular enumerations can be bound quite easily, for instance.

Simple query

Simple calculation:

```
// Select 10/5;
auto rows = storage.select(10/5);
```

This statement produces `std::vector<int>`.

```
// Select 10/5, 2*4;
auto rows = storage.select(columns(10/5, 2*4));
```

This statement produces `std::vector<std::tuple<int, int>>`

General SELECT Syntax:

```
SELECT DISTINCT column_list
FROM table_list
JOIN table ON join_condition
WHERE row_filter
ORDER BY column
LIMIT count OFFSET offset
GROUP BY column
HAVING group_filter;
```

One table select:

```
// SELECT name, id FROM User;
auto rows = storage.select(columns(&User::name, &User::id));
auto rows = storage.select(columns(&User::name, &User::id), from<User>());
```

These statements are equivalent and they produce a `std::vector<std::tuple<std::string, int>>`³. When the from clause is omitted there is an algorithm that detects all types present in a statement and adds all of them to the from clause. This works immediately when only one type is involved but sometimes we need to add joins to the other tables in which case it is best to use the explicit `from<>()`.

```
// SELECT * FROM User;
auto rows = storage.select(asterisk<User>()); // get all columns from User
Produces std::vector<std::tuple<std::string,int>>
```

```
auto objects = storage.get_all<User>(); // get all persistent instances of the User type
Produces std::vector<User>
```

³ Assuming the definition of columns as not nullable; otherwise it would be `std::vector<std::tuple<std::optional<std::string>,int>` if we use `std::optional` to create the name field

When dealing with large resultsets

We don't have to load whole result set into memory! We can iterate the collections!

```
for(auto& employee: storage.iterate<Employee>()) {  
    cout << storage.dump(employee) << endl;  
}  
  
for(auto& hero: storage.iterate<MarvelHero>(where(length(&MarvelHero::name) < 6))) {  
    cout << "hero = " << storage.dump(hero) << endl;  
}
```

Sorting rows

Order by

General syntax:

```
// SELECT select_list FROM table ORDER BY column_1 ASC, column_2 DESC;
```

Simple order by:

```
// SELECT "User"."first_name", "User"."last_name" FROM 'User' ORDER BY "User"."last_name" COLLATE NOCASE DESC
```

```
auto rows = storage.select(columns(&User::first_name, &User::last_name),  
                           order_by(&User::last_name).desc().collate_nocase());
```

```
// SELECT "User"."id", "User"."first_name", "User"."last_name" FROM 'User' ORDER BY "User"."last_name" DESC  
auto objects = storage.get_all<User>(order_by(&User::last_name).desc());
```

If desc() is omitted the ordering is ascending by default and if we omit the collation, binary is the default.

Compound order by:

```
auto rows = storage.select(columns(&User::name, &User::age), multi_order_by(  
    order_by(&User::name).asc(),
```

```

        order_by(&User::age).desc());
auto objects = storage.get_all<User>(multi_order_by(
    order_by(&User::name).asc(),
    order_by(&User::age).desc()));

```

Dynamic Order by

Sometimes the exact arguments that determine ordering is not known until at runtime, which is why we have this alternative:

```

auto orderBy = dynamic_order_by(storage);
orderBy.push_back(order_by(&User::firstName).asc());
orderBy.push_back(order_by(&User::id).desc());
auto rows = storage.get_all<User>(orderBy);

```

Ordering by a function

```

// SELECT ename, job from EMP order by substring(job, len(job)-1,2)
auto rows = storage.select(columns(&Employee::m_ename, &Employee::m_job),
    order_by(substr(&Employee::m_job, length(&Employee::m_job) - 1, 2)));

```

Dealing with NULLs when sorting⁴

```

// SELECT "Emp"."ename", "Emp"."salary", "Emp"."comm" FROM 'Emp' ORDER BY CASE WHEN "Emp"."comm" IS NULL THEN 0 ELSE
// 1 END DESC
auto rows = storage.select(
    columns(&Employee::m_ename, &Employee::m_salary, &Employee::m_commission),
    order_by(case_<int>()
        .when(is_null(&Employee::m_commission), then(0))
        .else_(1)
        .end()).desc());

```

This can of course be simplified like this below but using case_ is more powerful (e.g. when you have more than 2 values):

```

auto rows = storage.select(columns(&Employee::m_ename, &Employee::m_salary, &Employee::m_commission),
    order_by(is_null(&Employee::m_commission)).asc());

```

⁴ See CASE in this document

Sorting on a data dependent key

```
auto rows = storage.select(columns(&Employee::m_ename, &Employee::m_salary, &Employee::m_commission),
                             order_by(case_<double>()
                                         .when(is_equal(&Employee::m_job, "SalesMan"), then(&Employee::m_commission))
                                         .else_(&Employee::m_salary)
                                         .end()).desc());
```

Filtering Data

Select distinct

SELECT DISTINCT select_list FROM table;

```
// SELECT DISTINCT(name) FROM EMP5
auto names = storage.select(distinct(&Employee::name));
result is of type std::vector<std::optional<std::string>>
```

```
auto names = storage.select(distinct(columns(&Employee::name)));
result is of type std::vector<std::tuple<std::optional<std::string>>>
```

Where Clause

SELECT column_list FROM table WHERE search_condition;

Search condition can be formed from these clauses and their composition with and/or:

- WHERE column_1 = 100;
- WHERE column_2 IN (1,2,3);
- WHERE column_3 LIKE 'An%';
- WHERE column_4 BETWEEN 10 AND 20;
- WHERE expression1 Op expression2
 - Op can be any comparison operator:

⁵ Assume Employee::name is nullable

- = (== in C++)
- !=, <> (!= in C++)
- <
- >
- <=
- >=

```
// SELECT COMPANY.ID, COMPANY.NAME, COMPANY.AGE, DEPARTMENT.DEPT
// FROM COMPANY, DEPARTMENT
// WHERE COMPANY.ID = DEPARTMENT.EMP_ID;
auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(is_equal(&Employee::id, &Department::empId)));
```

```
auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(c(&Employee::id) == &Department::empId));
```

composite where clause: clause1 [and|or] clause2 ... and can also be &&, or can also be ||.

```
auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(c(&Employee::id) == &Department::empId) and c(&Employee::age) < 20);
```

```
auto objects = storage.get_all<User>(where(lesser_or_equal(&User::id, 2)
and (like(&User::name, "T%") or glob(&User::name, "*S"))))
```

the where clause can be also be used in UPDATE and DELETE statements.

Limit

Constrain the number of rows returned (limit) by a query optionally indicating how many rows to skip (offset).

```
// SELECT column_list FROM table LIMIT row_count;
auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(c(&Employee::id) == &Department::empId),
                           limit(4));
```

```
auto objects = storage.get_all<Employee>(limit(4));
```

```
// SELECT column_list FROM table LIMIT row_count OFFSET offset;
auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(c(&Employee::id) == &Department::empId),
                           limit(4, offset(3)));
```

```

auto objects = storage.get_all<Employee>(limit(4, offset(3)));

// SELECT column_list FROM table LIMIT offset, row_count;

auto rows = storage.select(columns(&Employee::id, &Employee::name, &Employee::age, &Department::dept),
                           where(c(&Employee::id) == &Department::empId),
                           limit(3, 4));

auto objects = storage.get_all<Employee>(limit(3, 4));

Using limit with order by:

// get the 2 employees with the second and third higher salary
auto rows = storage.select(columns(&Employee::name, &Employee::salary), order_by(&Employee::salary).desc(),
                           limit(2, offset(1)));

auto objects = storage.get_all<Employee>(order_by(&Employee::salary).desc(), limit(2, offset(1)));

```

Between Operator

Logical operator that tests whether a value is inside a range of values including the boundaries.

NOTE: BETWEEN can be used in the WHERE clause of the SELECT, DELETE, UPDATE and REPLACE statements.

```

// Syntax:
// test_expression BETWEEN low_expression AND high_expression

// SELECT DEPARTMENT_ID FROM departments
// WHERE manager_id
// BETWEEN 100 AND 200
auto rows = storage.select(&Department::id, where(between(&Department::managerId, 100, 200)));

// SELECT DEPARTMENT_ID FROM departments
// WHERE DEPARTMENT_NAME
// BETWEEN "D" AND "F"
auto rows = storage.select(&Department::id, where(between(&Department::dept, "D", "F")));

auto objects = storage.get_all<Department>(where(between(&Department::dept, "D", "F")));

```

In

Whether a value matches any value in a list or subquery, syntax being:
expression [NOT] IN (value_list|subquery);

```
// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id IN
//     (SELECT DEPARTMENT_ID FROM departments
//      WHERE location_id=1700);
auto rows = storage.select(columns(&Employee::firstName, &Employee::lastName, &Employee::departmentId),
                           where(in(&Employee::departmentId,
                                   select(&Department::id, where(c(&Department::locationId) == 1700))))));

// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id IN (10,20,30)
std::vector<int> ids{ 10,20,30 };
auto rows = storage.select(columns(&Employee::firstName, &Employee::departmentId),
                           where(in(&Employee::departmentId, ids)));

auto objects = storage.get_all<Employee>(where(in(&Employee::departmentId, {10,20,30} )));

// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id NOT IN (10,20,30)
std::vector<int> ids{ 10,20,30 };
auto rows = storage.select(columns(&Employee::firstName, &Employee::departmentId),
                           where(not_in(&Employee::departmentId, ids)));

auto objects = storage.get_all<Employee>(where(not_in(&Employee::departmentId, { 10,20,30 } )));

// SELECT "Emp"."empno", "Emp"."ename", "Emp"."job", "Emp"."salary", "Emp"."deptno" FROM 'Emp' WHERE
// (("Emp"."ename", "Emp"."job", "Emp"."salary") IN (
// SELECT "Emp"."ename", "Emp"."job", "Emp"."salary" FROM 'Emp'
// INTERSECT
// SELECT "Emp"."ename", "Emp"."job", "Emp"."salary" FROM 'Emp' WHERE (("Emp"."job" = "Clerk"))))
auto rows = storage.select(columns
(&Employee::m_empno, &Employee::m_ename, &Employee::m_job, &Employee::m_salary, &Employee::m_deptno),
                           where(c(std::make_tuple( &Employee::m_ename, &Employee::m_job, &Employee::m_salary))
```



```

.in(
select(intersect(
    select(columns(
        &Employee::m_ename, &Employee::m_job, &Employee::m_salary)),
    select(columns(
        &Employee::m_ename, &Employee::m_job, &Employee::m_salary),
        where(c(&Employee::m_job) == "Clerk")
)))));

```

which of course can be simplified to:

```

// SELECT "Emp"."empno", "Emp"."ename", "Emp"."job", "Emp"."salary", "Emp"."deptno" FROM 'Emp'
// WHERE(("Emp"."ename", "Emp"."job", "Emp"."salary")
// IN(SELECT "Emp"."ename", "Emp"."job", "Emp"."salary" FROM 'Emp' WHERE(("Emp"."job" = "Clerk"))))
auto rows = storage.select(columns(
    &Employee::m_empno, &Employee::m_ename, &Employee::m_job, &Employee::m_salary, &Employee::m_deptno),
    where(
        in(std::make_tuple(&Employee::m_ename, &Employee::m_job, &Employee::m_salary),
        select(columns(&Employee::m_ename, &Employee::m_job, &Employee::m_salary),
            where(c(&Employee::m_job) == "Clerk")))));

```

Like

Matches a pattern using 2 wildcards: % and _.

% matches 0 or more characters while _ matches any character. For characters in the ASCII range, the comparison is case insensitive; otherwise it is case sensitive.

```
SELECT column_list FROM table_name WHERE column_1 LIKE pattern;
```

```

auto whereCondition = where(like(&User::name, "S%"));
auto users = storage.get_all<User>(whereCondition);

auto rows = storage.select(&User::id, whereCondition);
auto rows = storage.select(like("ototo", "ot_to"));
auto rows = storage.select(like(&User::name, "%S%a"));
auto rows = storage.select(like(&User::name, "^%a").escape("^"));

```

Glob

Similar to the like operator but using UNIX wildcards like so:

- The asterisk (*) matches any number of characters (pattern Man* matches strings that start with Man)
- The question mark (?) matches exactly one character (pattern Man? matches strings that start with Man followed by any character)
- The list wildcard [] matches one character from the list inside the brackets. For instance [abc] matches either an a, a b or a c.
- The list wildcard can use ranges as in [a-zA-Z0-9]
- By using ^, we can match any character except those in the list ([^0-9] matches any non-numeric character).

```
auto rows = storage.select(columns(&Employee::lastName), where(glob(&Employee::lastName, "[^A-J]*")));
auto employees = storage.get_all<Employee>(where(glob(&Employee::lastName, "[^A-J]*")));
```

IS NULL

```
// SELECT
//     artists.ArtistId,
//     albumId
// FROM
//     artists
// LEFT JOIN albums ON albums.artistid = artists.artistid
// WHERE
//     albumid IS NULL;
auto rows = storage.select(columns(&Artist::artistId, &Album::albumId),
    left_join<Album>(on(c(&Album::artistId) == &Artist::artistId)),
    where(is_null(&Album::albumId)));
```

Dealing with NULL values in columns

```
// Transforming null values into real values
// SELECT COALESCE(comm,0), comm FROM EMP
auto rows = storage.select(columns(coalesce<double>(&Employee::m_commission, 0), &Employee::m_commission));
```

Joining tables

Table expressions are divided into join and nonjoin table expressions:

Table-expressions ::= join-table-expression | nonjoin-table-expression

Join-table-expression := table-reference CROSS JOIN table-reference

| table-reference [NATURAL] [join-type] JOIN table-reference [ON conditional-expression] | USING (column-commalist)]

| (join-table-expression)

SQLite join


In SQLite to query data from more than one table you can use INNER JOIN, LEFT JOIN or CROSS JOIN⁶. Each clause determines how rows from one table are “linked” to rows in another table. There is no explicit support for RIGHT JOIN or FULL OUTER JOIN. The expression OUTER is optional and does not alter the definition of the JOIN.

Cross join

Cross join is more accurately called the extended Cartesian product. If A and B are the tables from evaluation of the 2 table references then A CROSS JOIN B evaluates to a table consisting of all possible rows R such that R is the concatenation of a row from A and a row from B. In fact, the A CROSS JOIN B join expression is semantically equivalent to the following select-expression:

(SELECT A.* , B.* FROM A,B)

⁶ You can also use UNION or UNION ALL – see later in this document



(taken from [SQLite CROSS JOIN - w3resource](#))

Other joins

Table-reference [NATURAL] [join-type] JOIN table-reference
 [ON conditional-expression | using(column-commalist)]

Join type can be any of

- INNER⁷
- LEFT [OUTER]
- RIGHT [OUTER]
- FULL [OUTER]
- UNION⁸

With the following restrictions:

- NATURAL and UNION cannot both be specified
- If either NATURAL or UNION is specified, neither an ON clause nor a USING clause can be specified
- If neither NATURAL nor UNION is specified, then either an ON clause or a USING clause must be specified
- If join-type is omitted, INNER is assumed by default

It is important to realize that OUTER in LEFT, RIGHT and FULL has no effect on the overall semantics of the expression and is thus completely unnecessary. LEFT, RIGHT, FULL and UNION all have to do with NULLs so let's examine the other ones first:

1. Table-reference JOIN table-reference ON conditional-expression
2. Table-reference JOIN table-reference USING (column-commalist)
3. Table-reference NATURAL JOIN table-reference

Case 1 is equivalent to the following select-expression where cond is the conditional-expression:
(SELECT A.*, B.* FROM A,B WHERE cond)

In case 2, let the commalist of columns in the USING clause be unqualified C1, C2, ..., Cn, then it is equivalent to a case 1 with the following ON clause:
ON A.C1 = B.C1 AND A.C2 = B.C2 And ... A.Cn = B.Cn.

Finally case 3 is equivalent to case 2 where the USING clause contains all the columns that have the same names in A and B.

Joins having to do with NULLs (i.e. OUTER JOINS)

In the INNER joins, when we try to construct the ordinary join of 2 tables A and B, then any row that matches no row in the other table (under the relevant join condition) does not participate in the result. In an outer join such a row participates in the result: it appears exactly once, and the column positions that would have been filled with values from the other table, if such a mapping row had in fact existed, are filled with nulls instead. Therefore the outer join preserves nonmatching rows in the result whereas the inner join excludes them.

⁷ Default value if non specified (i.e., INNER JOIN is equivalent to JOIN)

⁸ We examine this concept later in this document

A LEFT OUTER JOIN of A and B, preserves rows from A with no matching rows from B. A RIGHT OUTER JOIN of A and B, preserves rows from B with no matching rows from A. A FULL OUTER JOIN preserves both. Lets analyze the particular cases for LEFT OUTER JOIN being that the other cases are similar:

We have three options in which to write our LEFT JOIN:

1. Table-reference LEFT JOIN table-reference ON conditional-expression
2. Table-reference LEFT JOIN table-reference USING (column-commalist)
3. Table-reference NATURAL LEFT JOIN table-reference

Case 1 can be represented as the following select statement:

```
SELECT A.*, B.* FROM A,B WHERE condition  
UNION ALL  
SELECT A.*, NULL, NULL, ...,NULL FROM A WHERE A.pkey NOT IN ( SELECT A.pkey FROM A,B WHERE condition)
```

Which means the UNION ALL of (a) the corresponding inner join and (b) the collection of rows excluded from the inner join, where there are as many NULL columns as there are columns in B.

For case 2, let the commalist of columns in the USING clause be C1, C2,..., Cn, all Ci unqualified and identifying a common column of A and B. Then the case becomes identical to a case 1 in which the condition has the form:

ON (A.C1 = B.C1 AND A.C2 = B.A2, ..., A.Cn = B.Cn)

For case 3, the commalist of colums to be used for case 2 is the collection of all common columns from A and B.

Example for Left Join:

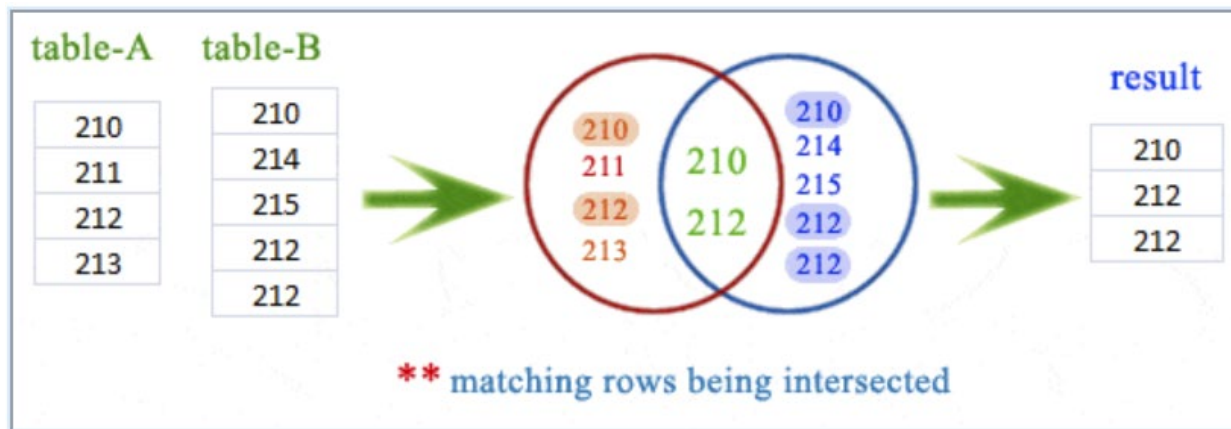
	AlbumId	Title	ArtistId	ArtistId:1	Name
1	1	Title	1	1	AC/DC
2	2	Title second	NULL	NULL	NULL

```
// SELECT
//     artists.ArtistId,
//     albumId
// FROM
//     artists
// LEFT JOIN albums ON albums.artistid = artists.artistid
// ORDER BY
//     albumid;
auto rows = storage.select(columns(&Artist::artistId, &Album::albumId),
    left_join<Album>(on(c(&Album::artistId) == &Artist::artistId)),
    order_by(&Album::albumId));
```

Example for Inner Join

```
// SELECT
//     trackid,
//     name,
//     title
// FROM
//     tracks
// INNER JOIN albums ON albums.albumid = tracks.albumid;
auto innerJoinRows0 = storage.select(columns(&Track::trackId, &Track::name, &Album::title),
    inner_join<Album>(on(
        c(&Track::albumId) == &Album::albumId)));
```

In this example, each row from tracks table is matched with a row from albums table according to the on clause. When this clause is true, then columns from the corresponding tables are displayed as an “extended row” – we are actually creating an anonymous type with attributes from the joined tables. The relationship between these tables is N tracks per 1 album. All N tracks with the same albumId are joined with the 1 album with matching columns as per the on clause.



Example for Natural Join

```
// SELECT doctor_id,doctor_name,degree,patient_name,vdate
// FROM doctors
// NATURAL JOIN visits
// WHERE doctors.degree="MD";
auto rows = storage.select(
    columns(&Doctor::doctor_id, &Doctor::doctor_name, &Doctor::degree, &Visit::patient_name, &Visit::vdate),
    natural_join<Visit>(),
    where(c(&Doctor::degree) == "MD"));

// SELECT doctor_id,doctor_name,degree,spl_descrip,patient_name,vdate
// FROM doctors
// NATURAL JOIN speciality
// NATURAL JOIN visits
// WHERE doctors.degree='MD';
auto rows = storage.select(columns(&Doctor::doctor_id,
                                &Doctor::doctor_name,
                                &Doctor::degree,
                                &Speciality::spl_descrip,
                                &Visit::patient_name,
                                &Visit::vdate),
    natural_join<Speciality>(),
    natural_join<Visit>(),
    where(c(&Doctor::degree) == "MD"));
```


Self join

```
// SELECT m.FirstName || ' ' || m.LastName,  
//         employees.FirstName || ' ' || employees.LastName  
// FROM employees  
// INNER JOIN employees m  
// ON m.ReportsTo = employees.EmployeeId  
using als = alias_m<Employee>;  
auto firstNames = storage.select(  
    columns(c(alias_column<als>(&Employee::firstName)) || " " || c(alias_column<als>(&Employee::lastName)),  
            c(&Employee::firstName) || " " || c(&Employee::lastName)),  
    inner_join<als>(on(alias_column<als>(&Employee::reportsTo) == c(&Employee::employeeId)))));
```

Full outer join

While SQLite does not support FULL OUTER JOIN, it is very easy to simulate it. Take these 2 classes/tables as an example, insert some data and do the “full outer join”:

```
struct Dog
{
    std::optional<std::string> type;
    std::optional<std::string> color;
};

struct Cat
{
    std::optional<std::string> type;
    std::optional<std::string> color;
};

using namespace sqlite_orm;

auto storage = make_storage(
    { "full_outer.sqlite" },
    make_table("Dogs", make_column("type", &Dog::type), make_column("color", &Dog::color)),
    make_table("Cats", make_column("type", &Cat::type), make_column("color", &Cat::color)));

storage.sync_schema();
storage.remove_all<Dog>();
storage.remove_all<Cat>();

storage.insert(into<Dog>(), columns(&Dog::type, &Dog::color), values(
    std::make_tuple("Hunting", "Black"), std::make_tuple("Guard", "Brown")));

storage.insert(into<Cat>(), columns(&Cat::type, &Cat::color), values(
    std::make_tuple("Indoor", "White"), std::make_tuple("Outdoor", "Black")));
```

```
// FULL OUTER JOIN simulation:
```

```
// SELECT d.type,  
// d.color,  
// c.type,  
// c.color  
// FROM dogs d  
// LEFT JOIN cats c USING(color)  
// UNION ALL  
// SELECT d.type,  
// d.color,  
// c.type,  
// c.color  
// FROM cats c  
// LEFT JOIN dogs d USING(color)  
// WHERE d.color IS NULL;
```

```
auto rows = storage.select(  
    union_all(select(columns(&Dog::type, &Dog::color, &Cat::type, &Cat::color),  
        left_join<Cat>(using_(&Cat::color))),  
        select(columns(&Dog::type, &Dog::color, &Cat::type, &Cat::color), from<Cat>(),  
            left_join<Dog>(using_(&Dog::color)), where(is_null(&Dog::color)))));
```

TYPE	COLOR	TYPE	COLOR
Hunting	Black	Outdoor	Black
Guard	Brown		
		Indoor	White

Grouping data

The group by clause is an optional clause of the select statement and enables us to take a selected group of rows into summary rows by values of one or more columns. It returns one row for each group and it is possible to apply an aggregate function such as MIN,MAX,SUM,COUNT or AVG – or one that you program yourself in `sqlite_orm`⁹!

The syntax is:

```
SELECT column_1, aggregate_function(column_2) FROM table GROUP BY column_1, column_2;
```

Group by

```
// If you want to know the total amount of salary on each customer, then GROUP BY query would be as follows:  
// SELECT NAME, SUM(SALARY)  
// FROM COMPANY  
// GROUP BY NAME;
```

```
auto salaryName = storage.select(columns(&Employee::name, sum(&Employee::salary)), group_by(&Employee::name));
```

Group by date example:

```
// SELECT (STRFTIME("%Y", "Invoices"."invoiceDate")) AS InvoiceYear,  
// (COUNT("Invoices"."id")) AS InvoiceCount FROM 'Invoices' GROUP BY InvoiceYear  
// ORDER BY InvoiceYear DESC  
struct InvoiceYearAlias : alias_tag {  
    static const std::string& get() {  
        static const std::string res = "INVOICE_YEAR";  
        return res;  
    }  
};
```

```
auto statement = storage.select(columns(as<InvoiceYearAlias>(strftime("%Y", &Invoice::invoiceDate)),  
as<InvoiceCountAlias>(count(&Invoice::id))), group_by(get<InvoiceYearAlias>()),  
order_by(get<InvoiceYearAlias>()).desc());
```

⁹ Take a peek at [create_scalar_function](#) and [create_aggregate_function](#) for details on how to define your custom functions in `sqlite_orm`

Having

While the where clause restricts the rows selected, the having clause selects data at the group level. For instance:


```
// SELECT NAME, SUM(SALARY)
// FROM COMPANY
// WHERE NAME is like "%l%"
// GROUP BY NAME
// HAVING SUM(SALARY) > 10000
auto namesWithHigherSalaries = storage.select(columns(&Employee::name, sum(&Employee::salary)),
                                              where(like(&Employee::name, "%l%")),
                                              group_by(&Employee::name).having(sum(&Employee::salary) > 10000));
```

Set operators

Union

The difference between UNION and JOIN is that the JOIN clause combines columns from multiple related tables while UNION combines rows from multiple similar tables. The UNION operator removes duplicate rows, whereas the UNION ALL operator does not. The rules for using UNION are as follows:

- Number of columns in all queries must be the same
- The corresponding columns must have compatible data types
- The column names of the first query determine the column names of the combined result set
- The group by and having clauses are applied to each individual query, not the final result set
- The order by apply to the combined result set, not within the individual result set



```
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// INNER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID
// UNION
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// LEFT OUTER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID;
auto rows = storage.select(
    union_(select(columns(&Department::employeeId, &Employee::name, &Department::dept),
        inner_join<Department>(on(is_equal(&Employee::id, &Department::employeeId))))),
    select(columns(&Department::employeeId, &Employee::name, &Department::dept),
        left_outer_join<Department>(on(is_equal(&Employee::id, &Department::employeeId))))));
```

Union all:

```
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// INNER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID
// UNION ALL
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// LEFT OUTER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID
auto rows = storage.select(
    union_all(select(columns(&Department::employeeId, &Employee::name, &Department::dept),
        inner_join<Department>(on(is_equal(&Employee::id, &Department::employeeId)))),
        select(columns(&Department::employeeId, &Employee::name, &Department::dept),
            left_outer_join<Department>(on(is_equal(&Employee::id, &Department::employeeId)))))));
```

Union ALL with order by:

```
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// INNER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID
// UNION ALL
// SELECT EMP_ID, NAME, DEPT
// FROM COMPANY
// LEFT OUTER JOIN DEPARTMENT
// ON COMPANY.ID = DEPARTMENT.EMP_ID
// ORDER BY NAME
auto rows = storage.select(
    union_all(select(columns(&Department::employeeId, &Employee::name, &Department::dept),
        inner_join<Department>(on(is_equal(&Employee::id, &Department::employeeId)))),
        select(columns(&Department::employeeId, &Employee::name, &Department::dept),
            left_outer_join<Department>(on(is_equal(&Employee::id, &Department::employeeId))),
            order_by(&Employee::name)))));
```

Stacking one resultset on top of another

```
// SELECT "Dept"."deptname" AS ENAME_AND_DNAME, "Dept"."deptno" FROM 'Dept'
// UNION ALL
// SELECT (QUOTE("-----")), NULL
// UNION ALL
// SELECT "Emp"."ename" AS ENAME_AND_DNAME, "Emp"."deptno" FROM 'Emp'
auto rows = storage.select(
    union_all(
        select(columns(as<NamesAlias>(&Department::m_deptname), as_optional(&Department::m_deptno))),
        select(union_all(
            select(columns(quote("-----"), std::optional<int>())),
            select(columns(as<NamesAlias>(&Employee::m_ename),
                as_optional(&Employee::m_deptno))))))));
```

Except

Compares the result sets of 2 queries and retains rows that are present only in the first result set. These are the rules:


- Number of columns in each query must be the same
- The order of the columns and their types must be comparable

Find all the dept_id in dept_master but not in emp_master:

```
// SELECT dept_id
// FROM dept_master
// EXCEPT
// SELECT dept_id
// FROM emp_master
auto rows = storage.select(except(select(&DeptMaster::deptId), select(&EmpMaster::deptId)));
```

Find all artists ids of artists who do not have any album in the albums table:

```
// SELECT ArtistId FROM artists EXCEPT SELECT ArtistId FROM albums;
auto rows = storage.select(except(select(&Artist::m_id), select(&Album::m_artist_id)));
```


Intersect

Compares the result sets of 2 queries and returns distinct rows that are output by both queries. Syntax:

SELECT select_list1 FROM table1 INTERSECT SELECT select_list2 FROM table2

These are the rules:

- Number of columns in each query must be the same
- The order of the columns and their types must be comparable



```
// SELECT dept_id
// FROM dept_master
// INTERSECT
// SELECT dept_id
// FROM emp_master
auto rows = storage.select(intersect(select(&DeptMaster::deptId), select(&EmpMaster::deptId)));
```

To find the customers who have invoices:

```
SELECT CustomerId FROM customers INTERSECT SELECT CustomerId FROM invoices ORDER BY CustomerId;
```

Subquery

Subquery

A subquery is a nested SELECT within another statement such as:

```
SELECT column_1 FROM table_1 WHERE column_1 = ( SELECT column_1 FROM table_2 );
```

```
// SELECT first_name, last_name, salary
// FROM employees
// WHERE salary >(
//     SELECT salary
//     FROM employees
//     WHERE first_name='Alexander');
auto rows = storage.select(
    columns(&Employee::firstName, &Employee::lastName, &Employee::salary),
    where(greater_than(&Employee::salary,
        select(&Employee::salary,
            where(is_equal(&Employee::firstName, "Alexander")))))));

// SELECT employee_id,first_name,last_name,salary
// FROM employees
// WHERE salary > (SELECT AVG(SALARY) FROM employees);
auto rows = storage.select(columns(
    &Employee::id, &Employee::firstName, &Employee::lastName, &Employee::salary),
    where(greater_than(&Employee::salary,
        select(avg(&Employee::salary)))));
```

```

// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id IN
//     (SELECT DEPARTMENT_ID FROM departments
//     WHERE location_id=1700);
auto rows = storage.select(
    columns(&Employee::firstName, &Employee::lastName, &Employee::departmentId),
    where(in(
        &Employee::departmentId,
        select(&Department::id, where(c(&Department::locationId) == 1700)))));

// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id IN (10,20,30)
std::vector<int> ids{ 10,20,30 };
auto rows = storage.select(columns(&Employee::firstName, &Employee::departmentId),
    where(in(&Employee::departmentId, ids)));

// SELECT * FROM employees
// WHERE department_id IN (10,20,30)
auto objects = storage.get_all<Employee>(where(in(&Employee::departmentId, {10,20,30} )));

// SELECT first_name, last_name, department_id
// FROM employees
// WHERE department_id NOT IN
//     (SELECT DEPARTMENT_ID FROM departments
//     WHERE manager_id
//     BETWEEN 100 AND 200);
auto rows = storage.select(
    columns(&Employee::firstName, &Employee::lastName, &Employee::departmentId),
    where(not_in(&Employee::departmentId,
        select(&Department::id, where(between(&Department::managerId, 100, 200))))));

// SELECT 'e'."LAST_NAME", 'e'."SALARY", 'e'."DEPARTMENT_ID" FROM 'employees' 'e'
// WHERE (('e'."SALARY" > (SELECT (AVG("employees"."SALARY")) FROM 'employees',
// 'employees' e WHERE (('employees"."DEPARTMENT_ID" = 'e'."DEPARTMENT_ID")))))
using als = alias_e<Employee>;
auto rows = storage.select(
    columns(alias_column<als>(&Employee::lastName),
    alias_column<als>(&Employee::salary),
    alias_column<als>(&Employee::departmentId)),

```

```

from<als>(),
where(greater_than(
alias_column<als>(&Employee::salary),
select(avg(&Employee::salary),
where(is_equal(&Employee::departmentId, alias_column<als>(&Employee::departmentId))))));

// SELECT first_name, last_name, employee_id, job_id
// FROM employees
// WHERE 1 <=
//      (SELECT COUNT(*) FROM Job_history
//      WHERE employee_id = employees.employee_id);
auto rows = storage.select(
    columns(&Employee::firstName, &Employee::lastName, &Employee::id, &Employee::jobId), from<Employee>(),
    where(lesser_or_equal(
        1,
        select(count<JobHistory>(), where(is_equal(&Employee::id, &JobHistory::employeeId))))));

```

```

SELECT albumid, title, (SELECT count(trackid) FROM tracks WHERE tracks.AlbumId = albums.AlbumId)    tracks_count FROM albums ORDER BY
tracks_count DESC;

```

Exists

Logical operator that checks whether subquery returns any rows. The subquery is a select statement that returns 0 or more rows. Syntax:

EXISTS (subquery)

```

// SELECT agent_code,agent_name,working_area,commission
// FROM agents
// WHERE exists
//      (SELECT *
//      FROM customer
//      WHERE grade=3 AND agents.agent_code=customer.agent_code)
// ORDER BY commission;
auto rows = storage.select(columns(&Agent::code, &Agent::name, &Agent::workingArea, &Agent::comission),
    from<Agent>(),
    where(exists(select(asterisk<Customer>(), from<Customer>(),
    where(is_equal(&Customer::grade, 3)
    and is_equal(&Agent::code, &Customer::agentCode))))),
    order_by(&Agent::comission));

```

```

// SELECT cust_code, cust_name, cust_city, grade
// FROM customer
// WHERE grade=2 AND EXISTS
//     (SELECT COUNT(*)
//      FROM customer
//      WHERE grade=2
//      GROUP BY grade
//      HAVING COUNT(*)>2);
auto rows = storage.select(columns(&Customer::code, &Customer::name, &Customer::city, &Customer::grade),
    where(is_equal(&Customer::grade, 2)
        and exists(select(count<Customer>(), where(is_equal(&Customer::grade, 2)),
            group_by(&Customer::grade),
            having(greater_than(count(), 2))))));

// SELECT "orders"."AGENT_CODE", "orders"."ORD_NUM", "orders"."ORD_AMOUNT", "orders"."CUST_CODE", 'c'."PAYMENT_AMT"
// FROM 'orders' INNER JOIN 'customer' 'c' ON('c'."AGENT_CODE" = "orders"."AGENT_CODE")
// WHERE(NOT(EXISTS
// (
//     SELECT 'd'."AGENT_CODE" FROM 'customer' 'd' WHERE((( 'c'."PAYMENT_AMT" = 7000) AND('d'."AGENT_CODE" =
//     'c'."AGENT_CODE")))))
// )
// ORDER BY 'c'."PAYMENT_AMT"

using als = alias_c<Customer>;
using als_2 = alias_d<Customer>;

double amount = 2000;
auto where_clause = select(alias_column<als_2>(&Customer::agentCode), from<als_2>(),
    where(is_equal(alias_column<als>(&Customer::paymentAmt), std::ref(amount)) and
        (alias_column<als_2>(&Customer::agentCode) == c(alias_column<als>(&Customer::agentCode)))));

amount = 7000;

auto rows = storage.select(columns(
    &Order::agentCode, &Order::num, &Order::amount, &Order::custCode, alias_column<als>(&Customer::paymentAmt)),
    from<Order>(),
    inner_join<als>(on(alias_column<als>(&Customer::agentCode) == c(&Order::agentCode))),
    where(not exists(where_clause)), order_by(alias_column<als>(&Customer::paymentAmt)));

```

More querying techniques

Case

We can add conditional logic to a query (an if else or switch statement in C++) by using the CASE expression. There are two syntaxes available and either can have column aliases (see below).

```
CASE case_expression
  WHEN case_expression = when_expression_1 THEN result_1
  WHEN case_expression = when_expression_2 THEN result_2
  ...
  [ ELSE result_else ]
END
```

```
// SELECT CASE "users"."country" WHEN "USA" THEN "Domestic" ELSE "Foreign" END
// FROM 'users' ORDER BY "users"."last_name" , "users"."first_name"
auto rows = storage.select(columns(
    case_<std::string>(&User::country)
    .when("USA", then("Domestic"))
    .else_("Foreign").end()),
    multi_order_by(order_by(&User::lastName), order_by(&User::firstName)));
```

```
CASE
  WHEN when_expression_1 THEN result_1
  WHEN when_expression_2 THEN result_2
  ...
  [ ELSE result_else ]
END
```

```

// SELECT ID, NAME, MARKS,
//     CASE
//     WHEN MARKS >=80 THEN 'A+'
//     WHEN MARKS >=70 THEN 'A'
//     WHEN MARKS >=60 THEN 'B'
//     WHEN MARKS >=50 THEN 'C'
//     ELSE 'Sorry!! Failed'
//     END
// FROM STUDENT;
auto rows = storage.select(columns(&Student::id,
                                   &Student::name,
                                   &Student::marks,
                                   case_<std::string>()
                                   .when(greater_or_equal(&Student::marks, 80), then("A+"))
                                   .when(greater_or_equal(&Student::marks, 70), then("A"))
                                   .when(greater_or_equal(&Student::marks, 60), then("B"))
                                   .when(greater_or_equal(&Student::marks, 50), then("C"))
                                   .else_("Sorry!! Failed")
                                   .end()));

```

Aliases for columns and tables

For tables:

```

// SELECT C.ID, C.NAME, C.AGE, D.DEPT
// FROM COMPANY AS C, DEPARTMENT AS D
// WHERE C.ID = D.EMP_ID;
using als_c = alias_c<Employee>;
using als_d = alias_d<Department>;
auto rowsWithTableAliases = storage.select(columns(
    alias_column<als_c>(&Employee::id),
    alias_column<als_c>(&Employee::name),
    alias_column<als_c>(&Employee::age),
    alias_column<als_d>(&Department::dept)),
    where(is_equal(alias_column<als_c>(&Employee::id), alias_column<als_d>(&Department::empId))));

```

For columns:

```
struct EmployeeIdAlias : alias_tag {
    static const std::string& get() {
        static const std::string res = "COMPANY_ID";
        return res;
    }
};

struct CompanyNameAlias : alias_tag {
    static const std::string& get() {
        static const std::string res = "COMPANY_NAME";
        return res;
    }
};

// SELECT COMPANY.ID as COMPANY_ID, COMPANY.NAME AS COMPANY_NAME, COMPANY.AGE, DEPARTMENT.DEPT
// FROM COMPANY, DEPARTMENT
// WHERE COMPANY_ID = DEPARTMENT.EMP_ID;
auto rowsWithColumnAliases = storage.select(columns(
    as<EmployeeIdAlias>(&Employee::id),
    as<CompanyNameAlias>(&Employee::name),
    &Employee::age,
    &Department::dept),
    where(is_equal(get<EmployeeIdAlias>(), &Department::empId))));
```

For columns and tables:

```
// SELECT C.ID AS COMPANY_ID, C.NAME AS COMPANY_NAME, C.AGE, D.DEPT
// FROM COMPANY AS C, DEPARTMENT AS D
// WHERE C.ID = D.EMP_ID;
auto rowsWithBothTableAndColumnAliases = storage.select(columns(
    as<EmployeeIdAlias>(alias_column<als_c>(&Employee::id)),
    as<CompanyNameAlias>(alias_column<als_c>(&Employee::name)),
    alias_column<als_c>(&Employee::age),
    alias_column<als_d>(&Department::dept)),
    where(is_equal(alias_column<als_c>(&Employee::id), alias_column<als_d>(&Department::empId))));
```


Applying aliases to CASE

```
struct GradeAlias : alias_tag {
    static const std::string& get() {
        static const std::string res = "Grade";
        return res;
    }
};

// SELECT ID, NAME, MARKS,
//     CASE
//     WHEN MARKS >=80 THEN 'A+'
//     WHEN MARKS >=70 THEN 'A'
//     WHEN MARKS >=60 THEN 'B'
//     WHEN MARKS >=50 THEN 'C'
//     ELSE 'Sorry!! Failed'
//     END as 'Grade'
//     FROM STUDENT;
auto rows = storage.select(columns(
    &Student::id,
    &Student::name,
    &Student::marks,
    as<GradeAlias>(case_<std::string>()
        .when(greater_or_equal(&Student::marks, 80), then("A+"))
        .when(greater_or_equal(&Student::marks, 70), then("A"))
        .when(greater_or_equal(&Student::marks, 60), then("B"))
        .when(greater_or_equal(&Student::marks, 50), then("C"))
        .else_("Sorry!! Failed")
        .end())));
```

Changing data

Inserting a single row into a table

```
INSERT INTO table (column1,column2 ,..) VALUES( value1, value2 ,...);
```

```
storage.insert(into<Invoice>(), columns(
    &Invoice::id, &Invoice::customerId, &Invoice::invoiceDate),
    values(std::make_tuple(1, 1, date("now"))));
```

Inserting an object

```
struct User {
    int id;           // primary key
    std::string name;
    std::vector<char> hash; // binary format
};
User alex{
    0,
    "Alex",
    {0x10, 0x20, 0x30, 0x40},
};
```

```
alex.id = storage.insert(alex); // inserts all non primary key columns, returns primary key when integral
```

Inserting several rows

```
storage.insert(into<Invoice>(),
    columns(&Invoice::id, &Invoice::customerId, &Invoice::invoiceDate),
    values(std::make_tuple(1, 1, date("now")),
        std::make_tuple(2, 1, date("now", "+1 year")),
        std::make_tuple(3, 1, date("now")),
        std::make_tuple(4, 1, date("now", "+1 year"))));
```

Inserting several objects via containers

If we want to insert or replace a group of persistent atoms, we can insert them into a container and provide iterators to the beginning and end of the desired range of objects, by means of the *insert_range* or *replace_range* methods of the storage type.

For instance:

```
std::vector<Department> des =
{
    Department{10, "Accounting", "New York"},
    Department{20, "Research", "Dallas"},
    Department{30, "Sales", "Chicago"},
    Department{40, "Operations", "Boston"}
};

std::vector<EmpBonus> bonuses =
{
    EmpBonus{-1, 7369, "14-Mar-2005", 1},
    EmpBonus{-1, 7900, "14-Mar-2005", 2},
    EmpBonus{-1, 7788, "14-Mar-2005", 3}
};

storage.replace_range(des.begin(), des.end());
storage.insert_range(bonuses.begin(), bonuses.end());
```

Recall that insert like statements do not set the primary keys while replace like statements copy all columns including primary keys. That should explain why we chose to replace the departments because they have explicit primary key values, and why we chose to insert the bonuses letting the database generate the primary key values.

Inserting several rows (becomes an update if primary key already exists)

```
// INSERT INTO COMPANY(ID, NAME, AGE, ADDRESS, SALARY)
// VALUES (3, 'Sofia', 26, 'Madrid', 15000.0)
//          (4, 'Doja', 26, 'LA', 25000.0)
// ON CONFLICT(ID) DO UPDATE SET NAME = excluded.NAME,
//                                AGE = excluded.AGE,
//                                ADDRESS = excluded.ADDRESS,
//                                SALARY = excluded.SALARY
storage.insert(
    into<Employee>(),
    columns(&Employee::id, &Employee::name, &Employee::age, &Employee::address, &Employee::salary),
    values(
        std::make_tuple(3, "Sofia", 26, "Madrid", 15000.0),
        std::make_tuple(4, "Doja", 26, "LA", 25000.0)),
    on_conflict(&Employee::id)
        .do_update(
            set(c(&Employee::name) = excluded(&Employee::name),
                c(&Employee::age) = excluded(&Employee::age),
                c(&Employee::address) = excluded(&Employee::address),
                c(&Employee::salary) = excluded(&Employee::salary)))));
```

Inserting only certain columns (provided the rest have either default_values, are nullable, are autoincrement or are generated):

```
// INSERT INTO Invoices("customerId") VALUES(2), (4), (8)
storage.insert(into<Invoice>(),
    columns(&Invoice::customerId),
    values(
        std::make_tuple(2),
        std::make_tuple(4),
        std::make_tuple(8)));

// INSERT INTO 'Invoices' ("customerId") VALUES (NULL)
Invoice inv{ -1, 1, std::nullopt };
storage.insert(inv, columns(&Invoice::customerId));
```

Inserting from select – getting rowid (since primary key is integral)

```
// INSERT INTO users SELECT "user_backup"."id", "user_backup"."name", "user_backup"."hash" FROM 'user_backup'  
  
storage.insert(into<User>(),  
    select(columns(&UserBackup::id, &UserBackup::name, &UserBackup::hash)))));  
auto r = storage.select(last_insert_rowid());
```

Inserting default values:

```
storage.insert(into<Artist>(), default_values());
```

Non-standard extension in SQLITE

Applies to UNIQUE, NOT NULL, CHECK and PRIMARY_KEY constraints, but not to FOREIGN KEY constraints.

For insert and update commands¹⁰, the syntax is INSERT OR Y or UPDATE OR Y where Y may be any of the following algorithms and the default conflict resolution algorithm is ABORT:

- ROLLBACK:
 - Aborts current statement with SQLITE_CONSTRAINT error and rolls back the current transaction; if no transaction active then behaves as ABORT
- ABORT
 - When constraint violation occurs returns with SQLITE_CONSTRAINT error and the current SQL statement backs out any changes made by it but changes caused by prior statements within the same transaction are preserved and the transaction remains active. This is the default conflict resolution algorithm.
- FAIL
 - Same as abort except that it does not back out prior changes of the current SQL statement... a foreign key constraint causes an ABORT
- IGNORE
 - Skips the one row that contains the constraint violation and continues processing subsequent rows of the SQL statement as if nothing went wrong: rows before and after the row with constraint violation are inserted or updated normally... a foreign key constraint violation causes an ABORT behavior
- REPLACE
 - When the constraint violation occurs of the UNIQUE or PRIMARY KEY type, the pre-existing rows causing the violation are deleted prior to inserting or updating the current row and the command continues executing normally. If a NOT NULL violation occurs, the NULL is replaced

¹⁰ The on conflict clause is used in the create table command with the same semantics

with the default value for that column if any exists, else the ABORT algorithm is used. For CHECK or foreign key violations, the algorithm works like ABORT. For the deleted rows, the delete triggers (if any) are fired if and only if recursive triggers¹¹ are enabled.

```
auto rows = storage.insert(or_abort(),
    into<User>(),
    columns(&User::id, &User::name),
    values(std::make_tuple(1, "The Weeknd")));

auto rows = storage.insert(or_fail(),
    into<User>(),
    columns(&User::id, &User::name),
    values(std::make_tuple(1, "The Weeknd")));

auto rows = storage.insert(or_ignore(),
    into<User>(),
    columns(&User::id, &User::name),
    values(std::make_tuple(1, "The Weeknd")));
auto rows = storage.insert(or_replace(),
    into<User>(),
    columns(&User::id, &User::name),
    values(std::make_tuple(1, "The Weeknd")));
auto rows = storage.insert(or_rollback(),
    into<User>(),
    columns(&User::id, &User::name),
    values(std::make_tuple(1, "The Weeknd")));
```

Update

This enables us to update data of existing rows in the table. The general syntax is like this:

```
UPDATE table SET column_1 = new_value_1, column_2 = new_value_2 WHERE search_condition;
```

¹¹ See PRAGMA recursive_triggers

Update several rows

```
// UPDATE COMPANY SET ADDRESS = 'Texas', SALARY = 20000.00 WHERE AGE < 30
storage.update_all(set(
    c(&Employee::address) = "Texas", c(&Employee::salary) = 20000.00),
    where(c(&Employee::age) < 30));

// UPDATE contacts
// SET phone = REPLACE12(phone, '410', '+1-410')
storage.update_all(set(
    c(&Contact::phone) = replace(&Contact::phone, "410", "+1-410")));
```

Update one row

```
// UPDATE products
// SET quantity = 5 WHERE id = 1;
storage.update_all(set(
    c(&Product::quantity) = 5),
    where(c(&Product::id) == 1));
```

Update an object

If student exists then update, else insert:

```
if(storage.count<Student>(where(c(&Student::id) == student.id))) {
    storage.update(student);
} else {
    studentId = storage.insert(student); // returns primary key
}

auto employee6 = storage.get<Employee>(6);

// UPDATE 'COMPANY' SET "NAME" = val1, "AGE" = val2, "ADDRESS" = "Texas" , "SALARY" = val4 WHERE "ID" = 6
employee6.address = "Texas";
storage.update(employee6); // actually this call updates all non-primary-key columns' values to passed object's
                          // fields
```

¹² See Core functions for definition of replace()

Delete Syntax

Since delete is a C++ keyword, remove and remove_all are used instead in sqlite_orm. The general syntax for DELETE is in SQL:

DELETE FROM table-name [WHERE expr]

Delete rows that satisfy a condition

```
// DELETE FROM artist WHERE artistname = 'Sammy Davis Jr.';
storage.remove_all<Artist>(
    where(c(&Artist::artistName) == "Sammy Davis Jr."));
```

Delete all objects of a certain type

```
// DELETE FROM Customer
storage.remove_all<Customer>();
```

Delete a certain object by giving its primary key

```
// DELETE FROM Customer WHERE id = 1;
storage.remove<Customer>(1);
```

Replace

If we want to set the primary key columns as well as the rest, we need to use replace instead of insert:

```
User john{
    2,
    "John",
    {0x10, 0x20, 0x30, 0x40},
};

// REPLACE INTO 'Users ("id", "name", "hash") VALUES (2, "John", {0x10, 0x20, 0x30, 0x40})
storage.replace(john);
```


Transactions

Transactions

SQLite is transactional in the sense that all changes and queries are atomic, consistent, isolated and durable, better known as ACID:

1. Atomic: the change cannot be broken into smaller ones: committing a transaction either applies every statement in it or none at all.
2. Consistent: the data must meet all validation rules before and after a transaction
3. Isolation: assume 2 transactions executing at the same time attempting to modify the same data. One of the 2 must wait until the other completes in order to maintain isolation
4. Durability: consider a transaction that commits but then the program crashes or the operating system crashes or there is a power failure to the computer. A transaction must ensure that the committed changes will persist even under such situations.

Sqlite has some pragmas that define exactly how these transactions are done and what level of durability they offer. For better durability less performance. Please see **PRAGMA schema.journal_mode** in [Pragma statements supported by SQLite...](#) and [Write-Ahead Logging \(sqlite.org\)](#) for detailed discussion.

NOTE: Changes to the database are faster if done within a transaction as in what follows:

```
storage.begin_transaction();
storage.replace(Employee{
    1,
    "Adams",
    "Andrew",
    "General Manager",
    {},
    "1962-02-18 00:00:00",
    "2002-08-14 00:00:00",
    "11120 Jasper Ave NW",
    "Edmonton",
    "AB",
    "Canada",
    "T5K 2N1",
    "+1 (780) 428-9482",
    "+1 (780) 428-3457",
    "andrew@chinookcorp.com",
});
storage.replace(Employee{
    2,
    "Edwards",
    "Nancy",
    "Sales Manager",
    std::make_unique<int>(1),
```

```

        "1958-12-08 00:00:00",
        "2002-05-01 00:00:00",
        "825 8 Ave SW",
        "Calgary",
        "AB",
        "Canada",
        "T2P 2T3",
        "+1 (403) 262-3443",
        "+1 (403) 262-3322",
        "nancy@chinookcorp.com",
    });
    storage.commit(); // or storage.rollback();
    storage.transaction([&storage] {
        storage.replace(Student{1, "Shweta", "shweta@gmail.com", 80});
        storage.replace(Student{2, "Yamini", "rani@gmail.com", 60});
        storage.replace(Student{3, "Sonal", "sonal@gmail.com", 50});
        return true; // commits
    });

```

NOTE: use of transaction guard implements RAII idiom

```

auto countBefore = storage.count<Object>();
try {
    auto guard = storage.transaction_guard();
    storage.insert(Object{0, "John"});
    storage.get<Object>(-1); // throws exception!
    REQUIRE(false);
} catch(...) {
    auto countNow = storage.count<Object>();
    REQUIRE(countBefore == countNow);
}

```

```

auto countBefore = storage.count<Object>();
try {
    auto guard = storage.transaction_guard();
    storage.insert(Object{0, "John"});
    guard.commit();
    storage.get<Object>(-1); // throws exception but transaction is not rolled back!
    REQUIRE(false);
} catch(...) {
    auto countNow = storage.count<Object>();
    REQUIRE(countNow == countBefore + 1);
}

```

Core functions

```

// SELECT name, LENGTH(name)
// FROM marvel
auto namesWithLengths = storage.select(
    columns(&MarvelHero::name,
           length(&MarvelHero::name))); // namesWithLengths is std::vector<std::tuple<std::string, int>>

// SELECT ABS(points)
// FROM marvel
auto absPoints = storage.select(
    abs(&MarvelHero::points)); // std::vector<std::unique_ptr<int>>
cout << "absPoints: ";
for(auto& value: absPoints)
{
    if(value) {
        cout << *value;
    } else {
        cout << "null";
    }
    cout << " ";
}
cout << endl;

// SELECT LOWER(name)
// FROM marvel
auto lowerNames = storage.select(
    lower(&MarvelHero::name));

```

```

// SELECT UPPER(abilities)
// FROM marvel
auto upperAbilities = storage.select(
    upper(&MarvelHero::abilities));

storage.transaction([&] {
    storage.remove_all<MarvelHero>();
    { // SELECT changes()

        auto rowsRemoved = storage.select(changes()).front();
        cout << "rowsRemoved = " << rowsRemoved << endl;
        assert(rowsRemoved == storage.changes());
    }
    { // SELECT total_changes()
        auto rowsRemoved = storage.select(total_changes()).front();
        cout << "rowsRemoved = " << rowsRemoved << endl;
        assert(rowsRemoved == storage.changes());
    }
    return false; // rollback
});
// SELECT CHAR(67, 72, 65, 82)
auto charString = storage.select(
    char_(67, 72, 65, 82)).front();
cout << "SELECT CHAR(67,72,65,82) = *" << charString << "*" << endl;

// SELECT LOWER(name) || '@marvel.com'
// FROM marvel
auto emails = storage.select(
    lower(&MarvelHero::name) || c("@marvel.com"));

// SELECT TRIM(' TechOnTheNet.com ')
auto string = storage.select(
    trim(" TechOnTheNet.com ")).front();

// SELECT TRIM('000123000', '0')
storage.select(
    trim("000123000", "0")).front()

```

```
// SELECT * FROM marvel ORDER BY RANDOM()
for(auto& hero: storage.iterate<MarvelHero>(order_by(sqlite_orm::random())) {
    cout << "hero = " << storage.dump(hero) << endl;
}
```

NOTE: Use iterate for large result sets because it does not load all the rows into memory

```
// SELECT ltrim(' TechOnTheNet.com is great!');
storage.select(ltrim(" TechOnTheNet.com is great!")).front();
```

Core functions can be used within prepared statements:

```
auto lTrimStatement = storage.prepare(select(
    ltrim("000123", "0")));

// SELECT ltrim('123123totn', '123');
get<0>(lTrimStatement) = "123123totn";
get<1>(lTrimStatement) = "123";
cout << "ltrim('123123totn', '123') = " << storage.execute(lTrimStatement).front() << endl;

// SELECT rtrim('TechOnTheNet.com ');
cout << "rtrim('TechOnTheNet.com ') = *" << storage.select(rtrim("TechOnTheNet.com ")).front() << "*" << endl;

// SELECT rtrim('123000', '0');
cout << "rtrim('123000', '0') = *" << storage.select(rtrim("123000", "0")).front() << "*" << endl;
// SELECT coalesce(NULL,20);
cout << "coalesce(NULL,20) = " << storage.select(coalesce<int>(std::nullopt, 20)).front() << endl;
cout << "coalesce(NULL,20) = " << storage.select(coalesce<int>(nullptr, 20)).front() << endl;

// SELECT substr('SQLite substr', 8);
cout << "substr('SQLite substr', 8) = " << storage.select(substr("SQLite substr", 8)).front() << endl;

// SELECT substr('SQLite substr', 1, 6);
cout << "substr('SQLite substr', 1, 6) = " << storage.select(substr("SQLite substr", 1, 6)).front() << endl;

// SELECT hex(67);
cout << "hex(67) = " << storage.select(hex(67)).front() << endl;

// SELECT quote('hi')
cout << "SELECT quote('hi') = " << storage.select(quote("hi")).front() << endl;
```

```

// SELECT hex(randomblob(10))
cout << "SELECT hex(randomblob(10)) = " << storage.select(hex(randomblob(10))).front() << endl;

// SELECT instr('something about it', 't')
cout << "SELECT instr('something about it', 't') = " << storage.select(instr("something about it", "t")).front();

struct o_pos : alias_tag {
    static const std::string& get() {
        static const std::string res = "o_pos";
        return res;
    }
};

// SELECT name, instr(abilities, 'o') o_pos
// FROM marvel
// WHERE o_pos > 0
auto rows = storage.select(columns(
    &MarvelHero::name, as<o_pos>(instr(&MarvelHero::abilities, "o"))),
    where(greater_than(get<o_pos>(), 0)));

// SELECT replace('AA B CC AAA','A','Z')
cout << "SELECT replace('AA B CC AAA','A','Z') = " << storage.select(replace("AA B CC AAA", "A", "Z")).front();

// SELECT replace('This is a cat','This','That')
cout << "SELECT replace('This is a cat','This','That') = "
    << storage.select(replace("This is a cat", "This", "That")).front() << endl;

// SELECT round(1929.236, 2)
cout << "SELECT round(1929.236, 2) = " << storage.select(round(1929.236, 2)).front() << endl;

// SELECT round(1929.236, 1)
cout << "SELECT round(1929.236, 1) = " << storage.select(round(1929.236, 1)).front() << endl;

// SELECT round(1929.236)
cout << "SELECT round(1929.236) = " << storage.select(round(1929.236)).front() << endl;

// SELECT unicode('A')
cout << "SELECT unicode('A') = " << storage.select(unicode("A")).front() << endl;

// SELECT typeof(1)
cout << "SELECT typeof(1) = " << storage.select(typeof_(1)).front() << endl;

```

```
// SELECT firstname, lastname, IFNULL(fax, 'Call:' || phone) fax
// FROM customers ORDER BY firstname
auto rows = storage.select(columns(
    &Customer::firstName, &Customer::lastName,
    ifnull<std::string>(&Customer::fax, "Call:" || c(&Customer::phone))),
    order_by(&Customer::firstName));
cout << "SELECT last_insert_rowid() = " << storage.select(last_insert_rowid()).front() << endl;
```

Data definition

Sqlite data types

SQLITE uses dynamic type system: the value stored in a column determines its data type, not the column's data type. You can even declare a column without specifying a data type. However columns created by `sqlite_orm` do have a declared data type.

SQLite provides primitive data types we call storage classes which are more general than a data type: INTEGER storage class includes 6 different types of integers.

Storage class	Meaning
NULL	NULL values mean missing information or unknown
Integer	Whole numbers with variable sizes such as 1,2,3,4 or 8 bytes
REAL	Real numbers with decimal values using 8 byte floats
TEXT	Stores character data of unlimited length. Supports various character encodings
BLOB	Binary large object that can store any kind of data of any length

The data type of a value is taken by these rules:

- If a literal has no enclosing quotes and decimal point or exponent, SQLite assigns the INTEGER storage class
- If a literal is inclosed by single or double quotes, SQLite assigns the TEXT storage class
- If a literal does not have quotes nor decimal points nor exponent, SQLite assigns the REAL storage class
- If a literal is NULL without quotes, it is assigned a NULL storage class
- If a literal has the format X'ABCD' or x'abcd' SQLite assigns BLOB storage class.
- Date and time can be stored as TEXT, INTEGER or REAL

How is data sorted when there are different storage classes?

Following these rules:

- NULL storage class has the lowest value... between NULL values there is no order
- The next higher storage classes are INTEGER and REAL, comparing them numerically
- The next higher storage class is TEXT, comparing them according to collation
- Highest storage class is BLOB, using the C function ***memcmp()*** to compare BLOB values

When using ORDER BY 2 steps are followed:

- Group values based on storage class: NULL, INTEGER, REAL, TEXT, BLOB
- Sort the values in each group

Therefore, even if the engine allows different types in one column, it is not a good idea!

Manifest typing and type affinity

- Manifest typing means that a data type is a property of a value stored in a column, not the property of the column in which the value is stored.. values of any type can be stored in a column
- Type affinity is the recommended type for data stored in that column – recommended, not required

```
SELECT typeof(100), typeof(10.0), typeof('100'), typeof(x'1000'), typeof(NULL);
```

In sqlite_orm typeof is typeof_.

Create table

In sqlite_orm we create the tables, indices, unique constraints, check constraints and triggers using the make_storage() function. Each data field of the struct we want to persist is mapped to one column in the table – but we don't have to add all of them: a struct may have non-storable fields. So first we define the types, normalize¹³ them and create the structs. These structs can be called “persistent atoms”.

```
struct User {
    int id;
    std::string name;
    std::vector<char> hash; // binary format
};

int main(int, char**) {
    using namespace sqlite_orm;
    auto storage = make_storage("blob.sqlite",
                               make_table("users",
                                           make_column("id", &User::id),
                                           make_column("name", &User::name, default_value("?")),
                                           make_column("hash", &User::hash)));
    storage.sync_schema();
}
```

This creates a database with name blob.sqlite and one table called users with 3 columns. The sync_schema() synchronizes the schema with the database but does not always work for existing tables. A workaround is to drop the tables and start the schema from zero. Adding uniqueness constraints to existing tables usually won't work... you need to version tables for doing some schema changes. It also defines a default value for column “name”.

¹³ As in relational normalization

CHECK constraint

Ensure values in columns meet specified conditions defined by an expression:

```
struct Contact {
    int id = 0;
    std::string firstName;
    std::string lastName;
    std::string email;
    std::string phone;
};

struct Product {
    int id = 0;
    std::string name;
    float listPrice = 0;
    float discount = 0;
};

auto storage = make_storage(":memory:",
    make_table("contacts",
        make_column("contact_id", &Contact::id),
        make_column("first_name", &Contact::firstName),
        make_column("last_name", &Contact::lastName),
        make_column("email", &Contact::email),
        make_column("phone", &Contact::phone),
        check(length(&Contact::phone) >= 10)),
    make_table("products",
        make_column("product_id", &Product::id, primary_key()),
        make_column("product_name", &Product::name),
        make_column("list_price", &Product::listPrice),
        make_column("discount", &Product::discount, default_value(0)),
        check(c(&Product::listPrice) >= &Product::discount and
            c(&Product::discount) >= 0 and c(&Product::listPrice) >= 0)));

storage.sync_schema();
```

This adds a check constraint and a column with default value.

Columns with specific collation and tables with Primary Key

```
struct User {
    int id;
    std::string name;
    time_t createdAt;
};

struct Foo {
    std::string text;
    int baz;
};

int main(int, char**) {
    using namespace sqlite_orm;
    auto storage = make_storage(
        "collate.sqlite",
        make_table("users",
            make_column("id", &User::id, primary_key()),
            make_column("name", &User::name),
            make_column("created_at", &User::createdAt)),
        make_table("foo", make_column("text", &Foo::text, collate_nocase()), make_column("baz", &Foo::baz)));
    storage.sync_schema();
}
```

This creates a case insensitive text column (other collations exist: collate_rtrim and collate_binary).

FOREIGN KEY constraint

```
struct Artist {
    int artistId;
    std::string artistName;
};

struct Track {
    int trackId;
    std::string trackName;
    std::optional<int> trackArtist; // must map to &Artist::artistId
};
```

```

int main(int, char** argv) {
    cout << "path = " << argv[0] << endl;

    using namespace sqlite_orm;
    { // simple case with foreign key to a single column without actions
        auto storage = make_storage("foreign_key.sqlite",
            make_table("artist",
                make_column("artistid", &Artist::artistId, primary_key()),
                make_column("artistname", &Artist::artistName)),
            make_table("track",
                make_column("trackid", &Track::trackId, primary_key()),
                make_column("trackname", &Track::trackName),
                make_column("trackartist", &Track::trackArtist),
                foreign_key(&Track::trackArtist).references(&Artist::artistId)));
        auto syncSchemaRes = storage.sync_schema();
        for (auto& p : syncSchemaRes) {
            cout << p.first << " " << p.second << endl;
        }
    }
}

```

This one defines a simple foreign key.

```

struct User {
    int id;
    std::string firstName;
    std::string lastName;
};

struct UserVisit {
    int userId;
    std::string userFirstName;
    time_t time;
};

```

```

int main() {
    using namespace sqlite_orm;

    auto storage = make_storage(
        {},
        make_table("users",
            make_column("id", &User::id, primary_key()),
            make_column("first_name", &User::firstName),
            make_column("last_name", &User::lastName),
            primary_key(&User::id, &User::firstName)),
        make_table("visits",
            make_column("user_id", &UserVisit::userId),
            make_column("user_first_name", &UserVisit::userFirstName),
            make_column("time", &UserVisit::time),
            foreign_key(&UserVisit::userId, &UserVisit::userFirstName)
                .references(&User::id, &User::firstName))),
    storage.sync_schema();
}

```

This defines a compound foreign key and a corresponding compound primary key.

AUTOINCREMENT property

```

struct DeptMaster {
    int deptId = 0;
    std::string deptName;
};

struct EmpMaster {
    int empId = 0;
    std::string firstName;
    std::string lastName;
    long salary;
    decltype(Department::deptId) deptId;
};

```

```

int main() {
    using namespace sqlite_orm;

    auto storage = make_storage("", // empty db name means in memory db
        make_table("dept_master",
            make_column("dept_id", &DeptMaster::deptId, primary_key(), autoincrement()),
            make_column("dept_name", &DeptMaster::deptName)),
        make_table("emp_master",
            make_column("emp_id", &EmpMaster::empId, autoincrement(), primary_key()),
            make_column("first_name", &EmpMaster::firstName),
            make_column("last_name", &EmpMaster::lastName),
            make_column("salary", &EmpMaster::salary),
            make_column("dept_id", &EmpMaster::deptId)));

    storage.sync_schema();
}

```

This defines both primary keys to be autoincrement(), so if you do not specify a value for the primary key one is created in a sequence. You may also determine the primary key explicitly using replace.

GENERATED COLUMNS

```

struct Product {
    int id = 0;
    std::string name;
    int quantity = 0;
    float price = 0;
    float totalValue = 0;
};

auto storage = make_storage({},
    make_table("products",
        make_column("id", &Product::id, primary_key()),
        make_column("name", &Product::name),
        make_column("quantity", &Product::quantity),
        make_column("price", &Product::price),
        make_column("total_value", &Product::totalValue,
            generated_always_as(&Product::price * c(&Product::quantity))));

storage.sync_schema();

```

This defines a generated column!

Databases may be created in memory if desired

By using the special name “:memory:” or just an empty name, SQLITE is instructed to create the database in memory.

```
struct RapArtist {
    int id;
    std::string name;
};

int main(int, char**) {
    auto storage = make_storage(":memory:",
                               make_table("rap_artists",
                                           make_column("id", &RapArtist::id, primary_key()),
                                           make_column("name", &RapArtist::name)));
    cout << "in memory db opened" << endl;
    storage.sync_schema();
}
```

This one is stored in memory (can also leave the dbname empty to achieve the same effect).

INDEX and UNIQUE INDEX

```
struct Contract {
    std::string firstName;
    std::string lastName;
    std::string email;
};

using namespace sqlite_orm;
```

```
// beware - put `make_index` before `make_table` cause `sync_schema` is called in reverse order
// otherwise you'll receive an exception
auto storage = make_storage(
    "index.sqlite",
    make_index("idx_contacts_name", &Contract::firstName, &Contract::lastName,
        where(length(&Contract::firstName) > 2)),
    make_unique_index("idx_contacts_email", indexed_column(&Contract::email).collate("BINARY").desc()),
    make_table("contacts",
        make_column("first_name", &Contract::firstName),
        make_column("last_name", &Contract::lastName),
        make_column("email", &Contract::email)));
```

This one allows you to create an index and a unique index.

DEFAULT VALUE for DATE columns

```
struct Invoice
{
    int id;
    int customerId;
    std::optional<std::string> invoiceDate;
};

using namespace sqlite_orm;

int main(int, char** argv) {
    cout << argv[0] << endl;

    auto storage = make_storage("aliases.sqlite",
        make_table("Invoices", make_column("id", &Invoice::id, primary_key(), autoincrement()),
            make_column("customerId", &Invoice::customerId),
            make_column("invoiceDate", &Invoice::invoiceDate, default_value(date("now")))));
```

this one defines the default value of invoiceDate to be the current date at the moment of insertion.

PERSISTENT collections

```
/**
 * This is just a mapped type.
 */
struct KeyValue {
    std::string key;
    std::string value;
};

auto& getStorage() {
    using namespace sqlite_orm;
    static auto storage = make_storage("key_value_example.sqlite",
                                     make_table("key_value",
                                                 make_column("key", &KeyValue::key, primary_key()),
                                                 make_column("value", &KeyValue::value)));
    return storage;
}

void setValue(const std::string& key, const std::string& value) {
    using namespace sqlite_orm;
    KeyValue kv{key, value};
    getStorage().replace(kv);
}

std::string getValue(const std::string& key) {
    using namespace sqlite_orm;
    if(auto kv = getStorage().get_pointer<KeyValue>(key)) {
        return kv->value;
    } else {
        return {};
    }
}
```

Implements a persistent map.

GETTERS and SETTERS

```
class Player {
    int id = 0;
    std::string name;
public:
    Player() {}
    Player(std::string name_) : name(std::move(name_)) {}
    Player(int id_, std::string name_) : id(id_), name(std::move(name_)) {}

    std::string getName() const {
        return this->name;
    }

    void setName(std::string name) {
        this->name = std::move(name);
    }

    int getId() const {
        return this->id;
    }

    void setId(int id) {
        this->id = id;
    }
};

int main(int, char**) {
    using namespace sqlite_orm;
    auto storage = make_storage("private.sqlite",
                               make_table("players",
                                           make_column("id",
                                                         &Player::setId, // setter
                                                         &Player::getId, // getter
                                                         primary_key()),
                                           make_column("name",
                                                         &Player::getName, // order between setter and getter doesn't matter.
                                                         &Player::setName)));

    storage.sync_schema();
}
```

This one uses getters and setters (note the order does not matter).

DEFINING THE SHEMA FOR SELF-JOINS

```
struct Employee {
    int employeeId;
    std::string lastName;
    std::string firstName;
    std::string title;
    std::unique_ptr<int> reportsTo;    // can also be std::optional<int> for nullable columns
    std::string birthDate;
    std::string hireDate;
    std::string address;
    std::string city;
    std::string state;
    std::string country;
    std::string postalCode;
    std::string phone;
    std::string fax;
    std::string email;
};

int main() {
    using namespace sqlite_orm;
    auto storage = make_storage("self_join.sqlite",
                               make_table("employees",
                                           make_column("EmployeeId", &Employee::employeeId, autoincrement(), primary_key()),
                                           make_column("LastName", &Employee::lastName),
                                           make_column("FirstName", &Employee::firstName),
                                           make_column("Title", &Employee::title),
                                           make_column("ReportsTo", &Employee::reportsTo),
                                           make_column("BirthDate", &Employee::birthDate),
                                           make_column("HireDate", &Employee::hireDate),
                                           make_column("Address", &Employee::address),
                                           make_column("City", &Employee::city),
                                           make_column("State", &Employee::state),
                                           make_column("Country", &Employee::country),
                                           make_column("PostalCode", &Employee::postalCode),
                                           make_column("Phone", &Employee::phone),
                                           make_column("Fax", &Employee::fax),
                                           make_column("Email", &Employee::email),
                                           foreign_key(&Employee::reportsTo).references(&Employee::employeeId)));
    storage.sync_schema();
}
```

SUBENTITIES

```
class Mark {
public:
    int value;
    int student_id;
};

class Student {
public:
    int id;
    std::string name;
    int roll_number;
    std::vector<decltype(Mark::value)> marks;
};

using namespace sqlite_orm;
auto storage = make_storage("subentities.sqlite",
    make_table("students",
        make_column("id", &Student::id, primary_key()),
        make_column("name", &Student::name),
        make_column("roll_no", &Student::roll_number)),
    make_table("marks",
        make_column("mark", &Mark::value),
        make_column("student_id", &Mark::student_id),
        foreign_key(&Mark::student_id).references(&Student::id)));
```

```

// inserts or updates student and does the same with marks
int addStudent(const Student& student) {
    auto studentId = student.id;
    if(storage.count<Student>(where(c(&Student::id) == student.id))) {
        storage.update(student);
    } else {
        studentId = storage.insert(student);
    }
    // insert all marks within a transaction
    storage.transaction([&] {
        storage.remove_all<Mark>(where(c(&Mark::student_id) == studentId));
        for(auto& mark: student.marks) {
            storage.insert(Mark{mark, studentId});
        }
        return true;
    });
    return studentId;
}

/**
 * To get student from db we have to execute two queries:
 * `SELECT * FROM students WHERE id = ?`
 * `SELECT mark FROM marks WHERE student_id = ?`
 */
Student getStudent(int studentId) {
    auto res = storage.get<Student>(studentId);
    res.marks = storage.select(&Mark::value, where(c(&Mark::student_id) == studentId));
    return res; // must be moved automatically by compiler
}

```

This one implements a sub-entity.

UNIQUENESS AT THE COLUMN AND TABLE LEVEL

```
struct Entry {
    int id;
    std::string uniqueColumn;
    std::unique_ptr<std::string> nullableColumn;
};

int main(int, char**) {
    using namespace sqlite_orm;
    auto storage = make_storage("unique.sqlite",
                               make_table("unique_test",
                                           make_column("id", &Entry::id, autoincrement(), primary_key()),
                                           make_column("unique_text", &Entry::uniqueColumn, unique()),
                                           make_column("nullable_text", &Entry::nullableColumn),
                                           unique(&Entry::id, &Entry::uniqueColumn)));
}
```

this one implements uniqueness at the column and table levels.

NOT NULL CONSTRAINT

Every data field of a persistent struct is by default not null. If we desire to allow nulls in a column, the type for the corresponding field must be one of these:

1. `Std::unique_ptr<T>`
2. `Std::shared_ptr<T>`
3. `Std::optional<T>`

VACUUM

Why do we need vacuum?

- Dropping database objects such as tables, views, indexes, or triggers marks them as free but the database size does not decrease.
- Every time you insert or delete data from tables, the index and tables become fragmented
- Insert, update and delete operations reduces the number of rows that can be stored in a single page => increases the number of pages necessary to hold a table => decreases cache performance and time to read/write
- Vacuum defragments the database objects, repacks individual pages ignoring the free spaces – it rebuilds the database and enables one to change database specific configuration parameters such as page size, page format and default encoding... just set new values using pragma and proceed with vacuum.

```
storage.vacuum();
```

Triggers

What is a Trigger?

A named database code that is executed automatically when an INSERT, UPDATE or DELETE statement is issued against the associated table.

Why do we need them?

- Auditing: log the changes in sensitive data (e.g. salary, email)
- To enforce complex business rules at the database level and prevent invalid transactions

Syntax:

```
CREATE TRIGGER [IF NOT EXISTS] trigger_name  
  [BEFORE|AFTER|INSTEAD OF14] [INSERT|UPDATE|DELETE]  
  ON table_name  
  [WHEN condition]  
BEGIN  
  statements;  
END;
```

Accessing old and new column values according to action

Action	Availability
INSERT	NEW is available
UPDATE	Both NEW and OLD are available
DELETE	OLD is available

¹⁴ Only allowed for views

Examples of Triggers

```
// CREATE TRIGGER validate_email_before_insert_leads
// BEFORE INSERT ON leads
// BEGIN
// SELECT
// CASE
// WHEN NEW.email NOT LIKE '%_@_%._%' THEN
// RAISE (ABORT,'Invalid email address')
// END;
// END;
make_trigger("validate_email_before_insert_leads",
    before()
    .insert()
    .on<Lead>()
    .begin(select(case_<int>()
        .when(not like(new_(&Lead::email), "%_@_%._%"),
            then(raise_abort("Invalid email address")))
        .end()))
    .end())
```

```
// CREATE TRIGGER log_contact_after_update
// AFTER UPDATE ON leads
// WHEN old.phone <> new.phone
// OR old.email <> new.email
// BEGIN
// INSERT INTO lead_logs (
//     old_id,
//     new_id,
//     old_phone,
//     new_phone,
//     old_email,
//     new_email,
//     user_action,
//     created_at
// )
// VALUES
// (
//     old.id,
//     new.id,
```

```

//      old.phone,
//      new.phone,
//      old.email,
//      new.email,
//      'UPDATE',
//      DATETIME('NOW')
//    ) ;
//  END;
make_trigger("log_contact_after_update",
  after()
  .update()
  .on<Lead>()
  .when(is_not_equal(old(&Lead::phone), new_(&Lead::phone)) and
    is_not_equal(old(&Lead::email), new_(&Lead::email)))
  .begin(insert(into<LeadLog>(),
    columns(&LeadLog::oldId,
      &LeadLog::newId,
      &LeadLog::oldPhone,
      &LeadLog::newPhone,
      &LeadLog::oldEmail,
      &LeadLog::newEmail,
      &LeadLog::userAction,
      &LeadLog::createdAt),
    values(std::make_tuple(
      old(&Lead::id),
      new_(&Lead::id),
      old(&Lead::phone),
      new_(&Lead::phone),
      old(&Lead::email),
      new_(&Lead::email),
      "UPDATE",
      datetime("NOW")))))
  .end())

```

```

// CREATE TRIGGER validate_fields_before_insert_fondos
//     BEFORE INSERT
//         ON Fondos
//         BEGIN
// SELECT CASE WHEN NEW.abrev = '' THEN RAISE(ABORT, "Fondo abreviacion empty") WHEN LENGTH(NEW.nombre) = 0 // THEN
RAISE(ABORT, "Fondo nombre empty") END;
// END;
make_trigger("validate_fields_before_insert_fondos",
    before()
    .insert()
    .on<Fondo>()
    .begin(select(case_<int>()
        .when(is_equal(new_(&Fondo::abreviacion), ""),
            then(raise_abort("Fondo abreviacion empty"))))
        .when(is_equal(length(new_(&Fondo::nombre)), 0),
            then(raise_abort("Fondo nombre empty"))))
        .end()))
    .end())

```

```

// CREATE TRIGGER validate_fields_before_update_fondos
//     BEFORE UPDATE
//         ON Fondos
//         BEGIN
// SELECT CASE WHEN NEW.abrev = '' THEN RAISE(ABORT, "Fondo abreviacion empty") WHEN LENGTH(NEW.nombre) = 0 // THEN
RAISE(ABORT, "Fondo nombre empty") END;
// END;
make_trigger("validate_fields_before_update_fondos",
    before()
    .update()
    .on<Fondo>()
    .begin(select(case_<int>()
        .when(is_equal(new_(&Fondo::abreviacion), ""),
            then(raise_abort("Fondo abreviacion empty"))))
        .when(is_equal(length(new_(&Fondo::nombre)), 0),
            then(raise_abort("Fondo nombre empty"))))
        .end()))
    .end())

```

Data migration

Sqlite_orm supports automatic schema migration but there are a few caveats. The method `storage.sync_schema(bool preserve)` tries to synchronize the on memory schema (called the storage schema) defined by the `make_storage` call, with the database schema. We are going to explore what this method can handle and what changes it takes care of and what changes it doesn't and what to do when it is not enough for our needs. First let's explore how it relates to the database:

1. Tables present in the database are not altered in any way nor dropped if they are not mentioned in the `make_table()` calls of `make_storage()` – therefore your C++ project is capable of dealing with a subset of the tables in a database if desired
2. Every table from storage is compared with its database analog and the following rules determine the outcome:
 - a. If table does not exist, it will be created
 - b. If table exists with excess columns the table will be dropped and recreated to match the columns defined in storage schema
 - i. The data in the table will be preserved regardless of the value of the parameter to `sync_schema`
 - c. If we add a column in the storage schema that is either nullable, is `generated_always_as` or has a default value, then:
 - i. The data in the table will be preserved regardless of the value of the parameter to `sync_schema`
 - d. If we add a column in the storage schema that is neither nullable nor has a default value, then:
 - i. The data in the table will be lost unless a backup is made because:
 1. The data cannot be loaded with current schema because the column does not yet exist in the database
 2. Unless the data in the table is test data, adding this kind of a column without an intermediate column ***is not reasonable***
 - e. If a table or a column in it differs from the storage schema in properties/constraints, the difference is sometimes detected, and the table will in those cases be dropped and recreated but when it is not detected, we need to force the drop/recreate manually as described below
 - i. Currently the differences that are detected include:
 1. Removing a primary key that exists in database
 - a. Data preserved: **you can remove a primary key and not lose any data**
 2. Adding a primary key that is absent from the database
 - a. Data lost regardless of preserve's value: **adding a primary key in storage will lose all data**
 - b. **How to preserve data in this case:**
 - i. See [How to drop data without losing it](#)
 3. Removing nullable in storage while it is nullable in database
 - a. Data lost regardless of preserve's value: **removing nullable in storage will lose all data**
 - b. **How to preserve data in this case:**
 - i. First, null values will be represented by 0 if type is integer or real and the empty string if text
 1. In case of a blob, will be represented as a vector of size 0.
 - ii. See [How to drop data without losing it](#)
 4. Adding nullable in storage while it is not nullable in database
 - a. Data lost regardless of preserve's value: **adding nullable in storage will lose all data**
 - b. **How to preserve data in this case:**

- i. First of all we need to decide if we are going to assume there are null values represented in the database by certain values, see [Interpret values in non-nullable column as nullable](#)
 - c. Once this is decided we make the usual sequence of statements:
 - i. See [How to drop data without losing it](#)
- 5. If we add a default value to the storage schema and none is present in the database, then
 - a. Data lost regardless of preserve's value: **adding default in storage will lose all data**
 - b. See [How to drop data without losing it](#)
- 6. If we remove a default value from the storage schema and it is present in the database, then
 - a. Data lost regardless of preserve's value: **removing default in storage will lose all data**
 - b. See [How to drop data without losing it](#)
- 7. If we add a generated_always_as to a column while it is not present in the database, then
 - a. Data preserved: **you can remove a primary key and not lose any data**
- 8. If we remove a generated_always_as from a column while it is present in the database, then
 - a. Data lost regardless of preserve's value: **removing generated_always_as from storage will lose all data**
 - b. See [How to drop data without losing it](#)
- ii. The differences that are not detectable (and thus not trigger a table drop/recreate) include:
 - 1. If we add a unique constraint to a column or a table, then
 - a. First, we must ensure that the existing data is unique and correct duplicates
 - i. See [Ensuring that a column contains unique values before making the column unique](#)
 - b. Table is not dropped, and constraint is not added to the database and no data is lost
 - c. We **must drop** and reload the table on purpose with these statements:
 - i. [How to drop data without losing it](#)
 - 2. If we remove a unique constraint to a column or a table, then
 - a. Table is not dropped, and constraint is not added to the database and no data is lost
 - b. We **must drop** and reload the table on purpose with these statements:
 - i. [How to drop data without losing it](#)
 - 3. If we add a check constraint which was not present in the database, then
 - a. Table is not dropped, and constraint is not added to the database and no data is lost
 - b. We **must drop** the table on purpose with these statements:
 - i. [How to drop data without losing it](#)
 - 4. If we remove a check constraint where one was present in the database, then
 - a. Do same steps as adding a check constraint
 - 5. If we add a foreign key constraint which wasn't present in the database, then
 - a. Table is not dropped, and constraint is not added to the database and no data is lost
 - b. Must force the dropping/recreate, see [How to drop data without losing it](#)
 - 6. If we remove a foreign key constraint which was present in the database, then
 - a. Table is not dropped, and constraint is not removed from the database and no data is lost
 - b. Must force the dropping/recreate, see [How to drop data without losing it](#)

- f. If a table is going to be dropped/recreated by `sync_schema()`, you must ensure that all tables with dependent rows are loaded, dropped in correct order, then call `sync_schema()`, and then reload in reverse order as explained in detail in
 - i. [How to drop data without losing it](#)

About correct order of dropping/loading tables

A table with dependent rows cannot be dropped – it would leave dangling references. Therefore, we must create a graph of all tables connected by edges from the table with a foreign key to a table referenced by that foreign key and we must load their data and then drop them from the ground up (i.e., starting at the leaf nodes). To load the data back into them, we need to choose an inverse order from that used in the dropping. This will ensure that the foreign keys point to existing primary keys, and we will avoid exceptions maybe leading to permanent data loss (unless we make persistent backups of the tables before we start the dropping). We could even make a backup copy of the entire database before attempting to call `sync_schema()`.

Example

Consider a simple schema with two tables, one for User and one for Job and suppose we want to drop and recreate them (or `sync_schema` is going to try doing this for us). We must be sure that tables with dependent rows do not exist when the referenced table is dropped, else we will get an exception.

```
static auto storage = make_storage(dbFilePath,
    make_unique_index("name_unique.idx", &User::name ),
    make_table("user_table",
        make_column("id", &User::id, primary_key()),
        make_column("name", &User::name),
        make_column("born", &User::born),
        make_column("job", &User::job),
        make_column("blob", &User::blob),
        make_column("age", &User::age ),
        check(length(&User::name) > 2),
        check(c(&User::born) != 0),
        foreign_key(&User::job).references(&Job::id)),
    make_table("job",
        make_column("id", &Job::id, primary_key(), autoincrement()),
        make_column("name", &Job::name, unique()),
        make_column("base_salary", &Job::base_salary)));
```

Now load and drop tables in correct order:

```
std::vector<User> users = storage.get_all<User>();  
storage.drop_table(storage.tablename<User>());  
std::vector<Job> jobs = storage.get_all<Job>();  
storage.drop_table(storage.tablename<Job>());
```

now call sync_schema() to propagate changes to database:

```
auto m = storage.sync_schema(true);    // we may inspect the return 'm' for information of actions performed
```

and reload tables in correct order (inverse of that used in dropping):

```
storage.replace_range(jobs.begin(), jobs.end());  
storage.replace_range(users.begin(), users.end());
```

How to drop data without losing it

When we have a difference between the storage schema and the database that is detectable by the sync_schema() function and it triggers table to be dropped and recreated or when the difference is not detectable we must do the following to keep schemas synchronized:

1. Call make_storage()
2. If the change would add a column that is not nullable and does not have a default value nor is generated, then consider creating an intermediate column with either property and then change the properties; *do not* add such a column in only one step because you will lose all the table data!
3. Load all data from the transitive dependent tables of the current table and the current table in the order specified by [About correct order of dropping/loading tables](#)¹⁵
4. If case is moving from a non-nullable to nullable column decide what (if at all) we are going to interpret as null values and modify the loaded data to change them to null values as described in [Interpret values in non-nullable column as nullable](#)
5. Drop tables in order given by point 2
6. Call sync_schema() – regardless of preserve value¹⁶
7. Replace all data into tables from the std::vectors in reverse order from that in which we dropped them

¹⁵ This will be possible in all cases except when adding a column with following attributes: not nullable, does not provide a default value, is not generated. This case should ALWAYS be avoided. Direct creation of such column types should not be done because we lose all data in the table. The workaround is: add the column desired but give it a nullable quality, or a default value or a generate value; then change the storage schema to remove the condition that allowed adding the column (i.e., nullable or default value or *generated value*)

¹⁶ We prefer to use a true value for this parameter

Interpret values in non-nullable column as nullable

If we are adding nullable to an existing column that is not nullable in the database, then we need to decide if we are going to interpret certain values as nullable and modify the vector's elements. For instance:

- For integer or real, is 0 to be taken as null?
- For text, is "" to be taken as null?
- For blob, is size() of std::vector<char> == 0 to be taken as null?
- If we decide these values should be treated as nulls, then we must transform the nullable column to std::nullopt following this pattern before replacing the vector into the table. For instance if the type of column job is integer or real, then we check whether its value is 0 and if so we replace it with std::nullopt which is interpreted as NULL in SQL:

```
std::transform(users.begin(), users.end(), users.begin(), [](User& user)
{
    if (user.job && user.job.value() == 0) { user.job = std::nullopt; }    return user;
});
```

Making a backup of the entire database

```
template<typename T>
void backup_db(T& storage, std::string db_name)
{
    namespace fs = std::filesystem;

    auto path_to_db_name = fs::path(db_name);
    auto stem = path_to_db_name.stem().string();
    auto backup_stem = stem + "_backup1.sqlite";
    auto backup_full_path = path_to_db_name.parent_path().append(backup_stem).string();
    storage.backup_to(backup_full_path);
}
```

Ensuring that a column contains unique values before making the column unique

If we have a persistent struct User with a column age which we want to declare unique, we must first detect if there are duplicates in the table.

Consider:

1. loading the table into a vector
2. sorting the table by age
3. find if there are repeated values by using adjacent_find algorithm
4. compare return value to end iterator of vector: if it is different then we have a duplicate which we must correct!

The code could be like this:

```
std::vector<User> users = storage.get_all<User>();
std::sort(users.begin(), users.end(), [](const User& l, const User& r) { return l.age < r.age; });
auto it = std::adjacent_find(users.begin(), users.end(), [](const User& l, const User& r) { return l.age == r.age; });
if( it!= users.end()) {
    // there are duplicates!
    User user = *it;           // points to duplicate
    auto age = user.age;       // duplicate age
}
```

SQLite tools

SQLiteStudio and Sqlite3 command shell

GUI open source full featured SQLite client downloadable from [SQLiteStudio](#), runs on Windows, Linux and MacOS X written in C++ using Qt 5.15.2 and SQLite 3.35.4.

Sqlite3.exe command shell and other command line utilities and even source code downloadable from [SQLite Download Page](#).

Sqlite_orm library and DSL downloadable from [fnc12/sqlite_orm at dev \(github.com\)](#). Do the following git command from a Visual Studio 2022 console:

```
Git clone --branch dev https://github.com/fnc12/sqlite\_orm.git sqlite_orm_dev_download_date
```

Where 'download_date' is the date you download it (e.g. March_1_2022).

Installing SQLite using vcpkg

Install Microsoft/vcpkg from [microsoft/vcpkg: C++ Library Manager for Windows, Linux, and MacOS \(github.com\)](https://github.com/microsoft/vcpkg) by following the [microsoft/vcpkg: C++ Library Manager for Windows, Linux, and MacOS \(github.com\)](https://github.com/microsoft/vcpkg). After installed, run at the command line the following:

```
> .\vcpkg\vcpkg install sqlite3:x64-windows
```

When you open Visual Studio 2022 the projects created will automatically find sqlite3.dll and sqlite3.lib.

SQLite import and export CSV

It is possible to import and export between comma separated texts and tables. This can be done with the command shell or with the GUI SQLiteStudio program (see [Import a CSV File Into an SQLite Table \(sqlitetutorial.net\)](https://sqlitetutorial.net/import-csv/) and [Export SQLite Database To a CSV File \(sqlitetutorial.net\)](https://sqlitetutorial.net/export-csv/)).

SQLite resources

[SQLite Resources \(sqlitetutorial.net\)](https://sqlitetutorial.net)

[SQLite Tutorial - An Easy Way to Master SQLite Fast](https://www.w3resource.com/sqlite-tutorial/)

[SQLite Home Page](https://www.w3resource.com/sqlite-tutorial/)

[SQLite Tutorial - w3resource](https://www.w3resource.com/sqlite-tutorial/)

[SQLite Exercises, Practice, Solution - w3resource](https://www.w3resource.com/sqlite-tutorial/)

[fnc12/sqlite_orm: ❤ SQLite ORM light header only library for modern C++ \(github.com\)](https://github.com/fnc12/sqlite-orm)

Debugging tips

For SELECT statements one can obtain the generated SQL by splitting the select command into a prepare + execute pair of commands. The prepare command produces a statement which can provide us with the SQL it represents and allows to make sure our query is as we have desired. For example:

```
auto statement = storage.prepare(select(columns(coalesce<double>(&Employee::m_commission, 0),
                                                &Employee::m_commission)));
auto sql = statement.expanded_sql();
auto rows = storage.execute(statement);
```

instead of the shortest:

```
auto rows = storage.select(columns(coalesce<double>(&Employee::m_commission, 0), &Employee::m_commission));
```

For information as to what storage.sync_schema() has done we can capture its return type which is a std::map of std::pair like so:

```
auto m = storage.sync_schema(true);
std::ostringstream oss;
for (auto& n : m) {
    oss << n.first << " " << n.second << "\t";
}
auto s = oss.str();
```

The Future of sqlite_orm

The most important features missing from sqlite_orm currently are support for **views** and **common table expressions**, in particular as represented by the WITH clause (see [The WITH Clause \(sqlite.org\)](#) and [The Simplest SQLite Common Table Expression Tutorial « Expensify Blog](#)).

An example of a dynamic from (which exemplify common table expressions) follows:

```
select depno, sum(salary) as total_sal, sum(bonus) as total_bonus from
(
    select e.empno,
           e.ename,
           e.salary,
           e.deptno,
           b.type,
           e.salary * case
             when b.type = 1 then .1
             when b.type = 2 then .2
             else .3
           end as bonus
    from emp e, emp_bonus b
   where e.empno = b.empno
        and e.deptno =20
) y
group by deptno
```

and an example of a WITH clause follows:

```
WITH RECURSIVE approvers(x) AS (
    SELECT 'Joanie'
    UNION ALL
    SELECT company.approver
    FROM company, approvers
    WHERE company.name=approvers.x AND company.approver IS NOT NULL
)
SELECT * FROM approvers;
```

References

[CPPTMP,2005] David Abrahams, Aleksey Gurtovoy. *C++ Template Metaprogramming*. Addison Wesley, 2005